

Due: 15MAR22

1. 20 Points

Load the CelebFaces Attributes (CelebA) dataset created by the Chinese University of Hong Kong (<http://mmlab.ie.cuhk.edu.hk/projects/CelebA.html>).

- a. Generate the "average" face using 1000 images in the dataset.
- b. Generate the average male image within this dataset.
- c. Generate the average female image within this dataset.
- d. Generate an image with random sampling.
- e. Why is the average image appear smooth while the image generated using random sampling grainy?

2. 40 Points

Using the MNIST dataset build a PixelCNN using a masked convolution with a 7x7 mask. Use ReLU and Sigmoid activation functions.

- a. Build a network with 7 residual blocks where there is no downsampling of the feature map and so padding is set to "same".
- b. Display the summary of the model.
- c. Compile the PixelCNN and run 10 epochs. Display the output of 25 images.
- d. Run for 25 epochs and display the output.
- e. Comment on the results of the PixelCNN along with computation speed.

3. 40 Points

Using the MNIST dataset build a DCGAN using a 7x7 mask. Use LeakyReLU as your activation function and with the padding set to "same". Use the Adam optimizer for both the generator and discriminator. (Code can be found at: <https://www.tensorflow.org/tutorials/g>)

- a. Generate a noise image with a normal distribution with pixel intensity values from 0 to 100 and display the image.
- b. Compile the model and run for 10 and 25 epochs and display the output.
- c. Comment on the computation speed and quality of images generated compared to PixelCNN.